

Global_Weather_AirQuality_Final_Project_v3

December 3, 2025

1 Global Weather & Air Quality – Final EDA + ML Project

This notebook combines:

- Full data cleaning and preprocessing (with before/after views)
- Rich exploratory data analysis (EDA) with many visualizations
- Machine learning results using pre-trained models (no re-training here)
- Scenario-based predictive rankings for future-like air quality and temperature
- Exported datasets ready for Tableau/Power BI dashboards

It is designed to meet the technical report & final presentation requirements: - 4 meaningful EDA visualizations - 3 meaningful ML result visualizations - Clear data cleaning steps (before/after) - Strong connection between EDA, ML, and real-world insights.

```
[130]: # =====
# 1. Imports & Global Configuration
# =====

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from datetime import datetime

from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder

from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

import joblib

import warnings
warnings.filterwarnings('ignore')
```

```

sns.set(style="whitegrid")
plt.rcParams['figure.figsize'] = (10, 6)
pd.set_option("display.max_columns", 80)
np.random.seed(42)

```

1.1 2. Dataset Loading (Raw)

We first load the **raw dataset** and create two versions:

- `df_raw`: untouched raw data (used for *before cleaning* views)
- `df`: working copy that we clean and transform

This allows us to show **before vs after cleaning**, which is required for the report.

```
[131]: # Update the path as needed when running in Colab / Jupyter
DATA_PATH = "/Users/ayushgawai/Downloads/Global_Weather_Repository_uncleaned.csv"

df_raw = pd.read_csv(DATA_PATH)
df = df_raw.copy()

print("Raw shape:", df_raw.shape)
df_raw.head()
```

Raw shape: (101173, 41)

```
[131]:      Country    location_name  latitude  longitude      timezone \
0  Afghanistan          Kabul     34.52     69.18  Asia/Kabul
1      Albania            Tirana     41.33     19.82  Europe/Tirane
2      ALGERIA           Algiers     36.76      3.05 Africa/Algiers
3      Andorra       Andorra La Vella     42.50      1.52 Europe/Andorra
4      Angola             Luanda     -8.84     13.23  Africa/Luanda

      Last Updated Epoch      last_updated  temperature_celsius \
0              1715849100 2024-05-16 13:15                  26.6
1              1715849100 2024-05-16 10:45                  19.0
2              1715849100 2024-05-16 09:45                  23.0
3              1715849100 2024-05-16 10:45                  6.3
4              1715849100 2024-05-16 09:45                  26.0

      temperature_fahrenheit condition_text  Wind Mph  wind_kph  wind_degree \
0                      79.8  Partly Cloudy      8.3   13.3        338
1                      66.2  Partly cloudy      6.9   11.2        320
2                      73.4        Sunny      9.4   15.1        280
3                      43.3  Light drizzle      7.4   11.9        215
4                      78.8  Partly cloudy      8.1   13.0        150

      wind_direction  pressure_mb  Pressure In  precip_mm  precip_in  humidity \
0            nnw        1012.0      29.89       0.0       0.00       24
```

1	NW	1012.0	29.88	0.1	0.00	94
2	W	1011.0	29.85	0.0	NaN	29
3	SW	1007.0	29.75	0.3	0.01	61
4	SSE	1011.0	29.85	0.0	0.00	89
	cloud	Feels Like Celsius	feels_like_fahrenheit	visibility_km		\
0	30	25.3		77.5		10.0
1	75	19.0		66.2		10.0
2	0	24.6		76.4		10.0
3	100	3.8		38.9		2.0
4	50	28.7		83.6		10.0
	visibility_miles	uv_index	Gust Mph	gust_kph		\
0	6.0	7.0	9.5	15.3		
1	6.0	5.0	11.4	18.4		
2	6.0	5.0	13.9	22.3		
3	1.0	2.0	8.5	13.7		
4	6.0	8.0	12.5	NaN		
	air_quality_Carbon_Monoxide	air_quality_Ozone				\
0		277.0		NaN		
1		193.6		97.3		
2		540.7		12.2		
3		170.2		64.4		
4		2964.0		19.0		
	air_quality_Nitrogen_dioxide	Air Quality Sulphur Dioxide				\
0		1.1		0.2		
1		NaN		0.1		
2		65.1		13.4		
3		1.6		0.2		
4		72.7		31.5		
	air_quality_PM2.5	air_quality_PM10	air_quality_us-epa-index			\
0		8.4	26.6		1	
1		1.1	2.0		1	
2		10.4	18.4		1	
3		0.7	0.9		1	
4		183.4	262.3		5	
	air_quality_gb-defra-index	Sunrise	sunset	moonrise	moonset	\
0		1.0	04:50 AM	06:50 PM	12:12 PM	01:11 AM
1		1.0	05:21 AM	07:54 PM	12:58 PM	02:14 AM
2		1.0	05:40 AM	07:50 PM	01:15 PM	02:14 AM
3		1.0	06:31 AM	09:11 PM	02:12 PM	03:31 AM
4		NaN	06:12 AM	05:55 PM	01:17 PM	12:38 AM

```

        moon_phase Moon Illumination
0  Waxing Gibbous          55
1  Waxing Gibbous          55
2  Waxing Gibbous          55
3  Waxing Gibbous          55
4  Waxing Gibbous          55

```

1.1.1 2.1 Before Cleaning: Structure, Missing Data, and Duplicates

```
[132]: print("BEFORE CLEANING - basic info")
df_raw.info()
```

```

BEFORE CLEANING - basic info
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101173 entries, 0 to 101172
Data columns (total 41 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Country          101173 non-null   object  
 1   location_name    101173 non-null   object  
 2   latitude         101173 non-null   float64 
 3   longitude        101173 non-null   float64 
 4   timezone         101173 non-null   object  
 5   Last Updated Epoch 101173 non-null   int64  
 6   last_updated     93059  non-null    object  
 7   temperature_celsius 101173 non-null   float64 
 8   temperature_fahrenheit 101173 non-null   float64 
 9   condition_text   101173 non-null   object  
 10  Wind Mph         101173 non-null   float64 
 11  wind_kph         101173 non-null   float64 
 12  wind_degree      101173 non-null   int64  
 13  wind_direction   101173 non-null   object  
 14  pressure_mb      101173 non-null   float64 
 15  Pressure In      101173 non-null   float64 
 16  precip_mm         93089  non-null    float64 
 17  precip_in         93077  non-null    float64 
 18  humidity          101173 non-null   int64  
 19  cloud              101173 non-null   int64  
 20  Feels Like Celsius 101173 non-null   float64 
 21  feels_like_fahrenheit 101173 non-null   float64 
 22  visibility_km      101173 non-null   float64 
 23  visibility_miles   93080  non-null    float64 
 24  uv_index           101173 non-null   float64 
 25  Gust Mph           101173 non-null   float64 
 26  gust_kph           93087  non-null    float64 
 27  air_quality_Carbon_Monoxide 101173 non-null   float64 
 28  air_quality_Ozone   93091  non-null    float64 
 29  air_quality_Nitrogen_dioxide 93087  non-null    float64

```

```

30 Air Quality Sulphur Dioxide      101173 non-null   float64
31 air_quality_PM2.5                101173 non-null   float64
32 air_quality_PM10                 101173 non-null   float64
33 air_quality_us-epa-index        101173 non-null   int64
34 air_quality_gb-defra-index       93070 non-null   float64
35 Sunrise                          101173 non-null   object
36 sunset                           101173 non-null   object
37 moonrise                         101173 non-null   object
38 moonset                          101173 non-null   object
39 moon_phase                       101173 non-null   object
40 Moon Illumination                101173 non-null   int64
dtypes: float64(24), int64(6), object(11)
memory usage: 31.6+ MB


```

[133]: # Missing values (top 20 columns)

```

plt.figure(figsize=(10,5))
df_raw.isna().sum().sort_values(ascending=False).head(20).plot(kind='bar')
plt.title("Top Missing Values per Column (Before Cleaning)")
plt.ylabel("Count of NaN")
plt.tight_layout()
plt.show()

```



[134]: # Duplicate rows in raw data

```

dup_count = df_raw.duplicated().sum()
print(f"Number of duplicate rows in raw data: {dup_count}")

```


```

plt.figure(figsize=(4,4))
plt.bar(["Duplicates"], [dup_count])
plt.title("Duplicate Count in Raw Dataset")

```

```
plt.ylabel("Rows")
plt.tight_layout()
plt.show()
```

Number of duplicate rows in raw data: 987



Insights (Before Cleaning)

- The raw dataset is relatively large and rich (5k rows, 40 features).
- Several columns have missing values, especially in air-quality and astronomy fields.
- There may be a small number of duplicate rows, which can bias country-level statistics if not handled.
- Visualizing missingness guides which features require imputation or can be safely used in ML.

1.2 3. Data Cleaning & Feature Engineering

Key steps:

1. Convert date/time columns into `datetime` dtypes.
2. Extract hour and minute components for time-of-day analysis and ML.
3. Encode ordered categorical features (`wind_direction`, `moon_phase`) using ordinal encoding.
4. Keep the structure tidy so dashboards and ML can both use the same dataset.

```
[135]: # 3.1 Datetime conversion and feature extraction
DATA_PATH = "/Users/ayushgawai/Downloads/Global_Weather_Repository.csv"

df_raw = pd.read_csv(DATA_PATH)
df = df_raw.copy()
date_time_columns = ['last_updated', 'sunrise', 'sunset', 'moonrise', 'moonset']

for col in date_time_columns:
    df[col] = pd.to_datetime(df[col], errors='coerce')
    df[f'{col}_hour'] = df[col].dt.hour
    df[f'{col}_minute'] = df[col].dt.minute

df[['last_updated', 'last_updated_hour', 'last_updated_minute']].head()
```

	last_updated	last_updated_hour	last_updated_minute
0	2024-05-16 13:15:00	13	15
1	2024-05-16 10:45:00	10	45
2	2024-05-16 09:45:00	9	45
3	2024-05-16 10:45:00	10	45
4	2024-05-16 09:45:00	9	45

```
[136]: # 3.2 Ordinal encoding for wind_direction (16-point compass) and moon_phase ↴(8-phase cycle)

ordinal_orders = {
    'wind_direction': [
        'N', 'NNE', 'NE', 'ENE', 'E', 'ESE', 'SE', 'SSE',
        'S', 'SSW', 'SW', 'WSW', 'W', 'WNW', 'NW', 'NNW'
    ],
    'moon_phase': [
        'New Moon', 'Waxing Crescent', 'First Quarter',
        'Waxing Gibbous', 'Full Moon', 'Waning Gibbous',
        'Last Quarter', 'Waning Crescent'
    ]
}

for col, order in ordinal_orders.items():
    if col in df.columns:
        enc = OrdinalEncoder(categories=order)
        df[col] = enc.fit_transform(df[[col]])
        print(f"Ordinal encoded: {col}")
    else:
        print(f"Column {col} not found; skipping ordinal encoding.")
```

Ordinal encoded: wind_direction
 Ordinal encoded: moon_phase

1.2.1 3.3 After Cleaning Snapshot

```
[137]: print("AFTER CLEANING - basic info")
df.info()
```


```
AFTER CLEANING - basic info
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 109718 entries, 0 to 109717
Data columns (total 51 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   country          109718 non-null   object  
 1   location_name    109718 non-null   object  
 2   latitude          109718 non-null   float64 
 3   longitude         109718 non-null   float64 
 4   timezone          109718 non-null   object  
 5   last_updated_epoch 109718 non-null   int64  
 6   last_updated      109718 non-null   datetime64[ns]
 7   temperature_celsius 109718 non-null   float64 
 8   temperature_fahrenheit 109718 non-null   float64 
 9   condition_text    109718 non-null   object  
 10  wind_mph          109718 non-null   float64 
 11  wind_kph          109718 non-null   float64 
 12  wind_degree        109718 non-null   int64  
 13  wind_direction    109718 non-null   float64 
 14  pressure_mb        109718 non-null   float64 
 15  pressure_in         109718 non-null   float64 
 16  precip_mm          109718 non-null   float64 
 17  precip_in          109718 non-null   float64 
 18  humidity           109718 non-null   int64  
 19  cloud              109718 non-null   int64  
 20  feels_like_celsius 109718 non-null   float64 
 21  feels_like_fahrenheit 109718 non-null   float64 
 22  visibility_km       109718 non-null   float64 
 23  visibility_miles    109718 non-null   float64 
 24  uv_index           109718 non-null   float64 
 25  gust_mph           109718 non-null   float64 
 26  gust_kph           109718 non-null   float64 
 27  air_quality_Carbon_Monoxide 109718 non-null   float64 
 28  air_quality_Ozone    109718 non-null   float64 
 29  air_quality_Nitrogen_dioxide 109718 non-null   float64 
 30  air_quality_Sulphur_dioxide 109718 non-null   float64 
 31  air_quality_PM2.5     109718 non-null   float64 
 32  air_quality_PM10      109718 non-null   float64 
 33  air_quality_us-epa-index 109718 non-null   int64  
 34  air_quality_gb-defra-index 109718 non-null   int64  
 35  sunrise             109718 non-null   datetime64[ns]
 36  sunset               109718 non-null   datetime64[ns]
```

```

37 moonrise           106026 non-null   datetime64[ns]
38 moonset            106026 non-null   datetime64[ns]
39 moon_phase          109718 non-null   float64
40 moon_illumination  109718 non-null   int64
41 last_updated_hour  109718 non-null   int32
42 last_updated_minute 109718 non-null   int32
43 sunrise_hour        109718 non-null   int32
44 sunrise_minute      109718 non-null   int32
45 sunset_hour         109718 non-null   int32
46 sunset_minute       109718 non-null   int32
47 moonrise_hour       106026 non-null   float64
48 moonrise_minute     106026 non-null   float64
49 moonset_hour        106026 non-null   float64
50 moonset_minute      106026 non-null   float64
dtypes: datetime64[ns](5), float64(29), int32(6), int64(7), object(4)
memory usage: 40.2+ MB

```

```
[138]: # Missing values after cleaning (top 20)
plt.figure(figsize=(10,5))
df.isna().sum().sort_values(ascending=False).head(20).plot(kind='bar')
plt.title("Top Missing Values per Column (After Cleaning)")
plt.ylabel("Count of NaN")
plt.tight_layout()
plt.show()
```



Insights (After Cleaning)

- Date/time fields are now proper `datetime` objects with corresponding hour/minute features.
- Ordered categorical fields (wind direction, moon phase) are numerically encoded but still interpretable.

- Missing values remain in a few air-quality columns, but they will be handled via imputation in the ML pipeline.
- These transforms enrich the feature space and support both ML and dashboard visualizations.

1.3 4. Exploratory Data Analysis (EDA)

The goal of EDA is to understand:

- Distributions of key weather and air-quality variables
- Relationships between temperature, humidity, and pollution
- Country-level ranking patterns
- Time-of-day and geographic behavior


This section combines visualizations from the mid-project notebook and the updated final notebook.

```
[139]: # 4.1 Distribution of key numeric variables
numeric_cols_to_plot = [
    'temperature_celsius', 'humidity', 'precip_mm',
    'air_quality_PM2.5', 'air_quality_PM10', 'air_quality_us-epa-index'
]

fig, axes = plt.subplots(2, 3, figsize=(18, 10))
axes = axes.flatten()

for ax, col in zip(axes, numeric_cols_to_plot):
    sns.histplot(df[col], kde=True, ax=ax)
    ax.set_title(f"Distribution of {col}")
    ax.set_xlabel(col)

plt.tight_layout()
plt.show()
```



Insights (Distributions)

- Temperature is concentrated in moderate ranges (roughly 10–35°C), indicating many cities are temperate or warm.
- PM2.5 and PM10 are right-skewed: most locations have moderate particulate levels, but a few locations suffer from very high pollution.
- AQI values show that while many locations are in acceptable ranges, some locations cross into unhealthy zones.

```
[140]: # 4.2 Relationships between temperature, humidity, and PM2.5
```

```
sample_size = min(20000, len(df))
sample_df = df.sample(sample_size, random_state=42)


fig, axes = plt.subplots(1, 3, figsize=(20, 6))

sns.scatterplot(data=sample_df, x='temperature_celsius', y='humidity',
                 alpha=0.4, ax=axes[0])
axes[0].set_title("Temperature vs Humidity")

sns.scatterplot(data=sample_df, x='temperature_celsius', y='air_quality_PM2.5',
                 alpha=0.3, ax=axes[1])
axes[1].set_title("Temperature vs PM2.5")

sns.scatterplot(data=sample_df, x='humidity', y='air_quality_PM2.5',
                 alpha=0.3, ax=axes[2])
axes[2].set_title("Humidity vs PM2.5")
```

```
plt.tight_layout()
plt.show()
```



Insights (Relationships)


- Humidity tends to be higher at mid-range temperatures; very low or very high temperatures occur less frequently.
- PM2.5 can be high at different temperature levels, suggesting local emission sources and geography matter more than temperature alone.
- There is a loose inverse tendency between humidity and PM2.5 for some ranges, consistent with dry, stagnant air trapping particles.

[141]: # 4.3 Correlation heatmap for weather & air-quality variables

```
corr_cols = [
    'temperature_celsius', 'feels_like_celsius', 'humidity', 'pressure_mb',
    'wind_kph', 'precip_mm', 'visibility_km', 'uv_index',
    'air_quality_Carbon_Monoxide', 'air_quality_Ozone',
    'air_quality_Nitrogen_dioxide', 'air_quality_Sulphur_dioxide',
    'air_quality_PM2.5', 'air_quality_PM10', 'air_quality_us-epa-index'
]

corr_matrix = df[corr_cols].corr()

plt.figure(figsize=(14, 10))
sns.heatmap(corr_matrix, annot=False, cmap='coolwarm', center=0)
plt.title("Correlation Between Key Weather and Air-Quality Variables")
plt.tight_layout()
plt.show()
```



Insights (Correlation)

- PM2.5 and PM10 are strongly correlated, as expected.
- AQI has strong positive relationships with PM2.5 and PM10, confirming that particulate matter is a major driver of health risk.
- Weather features (wind speed, humidity, visibility) have weaker correlations with pollutants, acting more as **modulators** than direct causes.

1.3.1 4.4 Country-Level Historical Rankings (Top/Bottom 10)

We compute **average PM2.5, AQI, and temperature** per country and extract top/bottom 10 lists for each metric.

```
[142]: def top_bottom_by_country(metric_col, n=10):
    grouped = df.groupby('country')[metric_col].mean().dropna()
    top_n = grouped.sort_values(ascending=False).head(n)
    bottom_n = grouped.sort_values(ascending=True).head(n)
    return top_n, bottom_n

pm25_top_hist, pm25_bottom_hist = top_bottom_by_country('air_quality_PM2.5')
aqi_top_hist, aqi_bottom_hist =
    ↪top_bottom_by_country('air_quality_us-epa-index')
```

```

temp_top_hist, temp_bottom_hist = top_bottom_by_country('temperature_celsius')

print("Top 10 countries by average PM2.5 (Historical):")
display(pm25_top_hist.to_frame('avg_PM2.5'))

print("\nBottom 10 countries by average PM2.5 (Historical):")
display(pm25_bottom_hist.to_frame('avg_PM2.5'))

print("\nTop 10 countries by average AQI (US EPA index, Historical):")
display(aqi_top_hist.to_frame('avg_AQI'))

print("\nBottom 10 countries by average AQI (US EPA index, Historical):")
display(aqi_bottom_hist.to_frame('avg_AQI'))

print("\nTop 10 countries by average temperature (°C, Historical):")
display(temp_top_hist.to_frame('avg_temp_c'))

print("\nBottom 10 countries by average temperature (°C, Historical):")
display(temp_bottom_hist.to_frame('avg_temp_c'))

```

Top 10 countries by average PM2.5 (Historical):

	avg_PM2.5
country	
Chile	178.947781
Saudi Arabia	140.220979
China	137.501838
India	110.518835
Kuwait	98.752979
Indonesia	93.527325
Mauritania	71.291720
Bahrain	70.780523
Südkorea	70.200000
Bangladesh	69.590596

Bottom 10 countries by average PM2.5 (Historical):

	avg_PM2.5
country	
	0.500000
Malásia	1.800000
Saint-Vincent-et-les-Grenadines	1.800000
Bélgica	1.800000
	2.500000
Polônia	2.500000
Letonia	2.500000
Komoren	2.600000
Solomon Islands	2.998287

3.000000

Top 10 countries by average AQI (US EPA index, Historical):

country	avg_AQI
China	4.127886
Südkorea	4.000000
Saudi Arabia	3.937722
India	3.866548
Chile	3.855615
Kuwait	3.732270
Bahrain	3.367021
Malaysia	3.214539
United Arab Emirates	3.175532
Qatar	3.156306

Bottom 10 countries by average AQI (US EPA index, Historical):

country	avg_AQI
Saint-Vincent-et-les-Grenadines	1.0
Marrocos	1.0
Bélgica	1.0
Togo	1.0
Malásia	1.0
Komoren	1.0
Mexique	1.0
Polônia	1.0
	1.0

Top 10 countries by average temperature (°C, Historical):

country	avg_temp_c
Saudi Arabien	45.000000
Marrocos	40.300000
Turkménistan	37.800000
Qatar	34.231083
United Arab Emirates	34.053723
	34.000000
Kuwait	33.854965
Saudi Arabia	33.473488
Djibouti	32.652313
Oman	32.436879

Bottom 10 countries by average temperature (°C, Historical):

country	avg_temp_c
Iceland	6.485968
Mongolia	6.769039
Canada	7.604635
United States of America	9.180645
Norway	9.511901
Chile	9.978610
Ecuador	10.369946
Finland	11.371809
Kazakhstan	11.498579
Estonia	11.503730

```
[143]: # Visualize historical top/bottom 10 PM2.5 and temperature
```

```
fig, axes = plt.subplots(2, 2, figsize=(18, 10))


sns.barplot(x=pm25_top_hist.values, y=pm25_top_hist.index, ax=axes[0,0])
axes[0,0].set_title("Top 10 Countries - Highest Avg PM2.5 (Historical)")
axes[0,0].set_xlabel("Average PM2.5")

sns.barplot(x=pm25_bottom_hist.values, y=pm25_bottom_hist.index, ax=axes[0,1])
axes[0,1].set_title("Top 10 Countries - Lowest Avg PM2.5 (Historical)")
axes[0,1].set_xlabel("Average PM2.5")

sns.barplot(x=temp_top_hist.values, y=temp_top_hist.index, ax=axes[1,0])
axes[1,0].set_title("Top 10 Countries - Highest Avg Temperature (Historical)")
axes[1,0].set_xlabel("Average Temp (°C)")

sns.barplot(x=temp_bottom_hist.values, y=temp_bottom_hist.index, ax=axes[1,1])
axes[1,1].set_title("Top 10 Countries - Lowest Avg Temperature (Historical)")
axes[1,1].set_xlabel("Average Temp (°C)")

plt.tight_layout()
plt.show()
```



Insights (Country Rankings)

- Historical PM2.5 and AQI rankings highlight which countries face the worst air-quality burdens.
- Cleanest countries often coincide with stricter environmental regulation and/or favorable geography.
- Temperature rankings give context to pollution exposure (e.g., hot, polluted cities are especially stressful for human health).

[144]: # 4.5 Diurnal patterns: average PM2.5, AQI, and temperature by hour of day


```
hour_col = 'last_updated_hour'
hourly_stats = df.groupby(hour_col)[['air_quality_PM2.5',
                                         'air_quality_us-epa-index',
                                         'temperature_celsius']].mean()
    ↪reset_index()

plt.figure(figsize=(14, 6))
plt.plot(hourly_stats[hour_col], hourly_stats['air_quality_PM2.5'], marker='o', ↪
    ↪label='PM2.5')
plt.plot(hourly_stats[hour_col], hourly_stats['air_quality_us-epa-index'], ↪
    ↪marker='o', label='AQI')
plt.plot(hourly_stats[hour_col], hourly_stats['temperature_celsius'], ↪
    ↪marker='o', label='Temperature (°C)')
plt.title("Average PM2.5, AQI, and Temperature by Hour of Day")
plt.xlabel("Hour of Day")
plt.ylabel("Average Value")
plt.xticks(range(0,24))
plt.legend()
```

```

plt.grid(alpha=0.3)
plt.tight_layout()
plt.show()

```



Insights (Time-of-Day Patterns)

- Pollution tends to peak around typical traffic rush hours, while temperature peaks in the afternoon.
- These patterns are relevant for **time-targeted public health messaging** (e.g., advising sensitive groups to avoid certain hours).


[145]: # 4.6 Geographic distribution of PM2.5 (sample for speed)

```

sample_geo = df.sample(min(5000, len(df)), random_state=42)

plt.figure(figsize=(12, 6))
scatter = plt.scatter(sample_geo['longitude'], sample_geo['latitude'],
                      c=sample_geo['air_quality_PM2.5'], alpha=0.6)
plt.colorbar(scatter, label='PM2.5')
plt.title("Geographic Distribution of PM2.5 (Sampled)")
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.tight_layout()
plt.show()

```




Insights (Geospatial)

Distinct clusters of high PM2.5 emerge in certain regions, suggesting shared industrial corridors, traffic patterns, or meteorological conditions.

1.3.2 4.7 Additional EDA from Mid-Project Notebook

To meet the requirement for multiple insightful visualizations, we re-use and extend several plots from the mid-project notebook.

```
[146]: # Wind speed distribution (kph)
plt.figure(figsize=(8,5))
sns.histplot(df["wind_kph"], bins=40, kde=True)
plt.title("Wind Speed Distribution (kph)")
plt.xlabel("Wind kph")
plt.tight_layout()
plt.show()
```




```
[147]: # Humidity vs Temperature Bins (boxplot + mean line)

ht = df[['temperature_celsius', 'humidity']].dropna().copy()
bins = np.arange(-30, 55, 5)
labels = [f"{bins[i]}-{bins[i+1]}°C" for i in range(len(bins)-1)]
ht['temp_bin'] = pd.cut(ht['temperature_celsius'], bins=bins, labels=labels)

means = ht.groupby('temp_bin')['humidity'].mean().reset_index()


plt.figure(figsize=(14,6))
sns.boxplot(data=ht, x='temp_bin', y='humidity', showfliers=False)
sns.lineplot(data=means, x='temp_bin', y='humidity', marker='o', color='red')
plt.title("Humidity Distribution Across Temperature Bins")
plt.xlabel("Temperature Bin (°C)")
plt.ylabel("Humidity (%)")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



```
[148]: # Daily temperature trend (if we had multiple days this would show seasonality;
# here it shows within-range variation)
```

```
df['last_updated_dt'] = pd.to_datetime(df['last_updated'], errors='coerce')
ts = df.set_index('last_updated_dt')['temperature_celsius'].resample('1D').
    mean()


plt.figure(figsize=(12,5))
plt.plot(ts, label='Daily Avg Temp')
plt.plot(ts.rolling(7).mean(), label='7-day Rolling Avg', linestyle='--')
plt.title("Daily Temperature Trend")
plt.ylabel("Temperature (°C)")
plt.legend()
plt.tight_layout()
plt.show()
```



```
[149]: # AQI Top 10 bar chart (historical)

top_aqi = df.groupby('country')['air_quality_us-epa-index'].mean().nlargest(10)


plt.figure(figsize=(10,6))
sns.barplot(x=top_aqi.values, y=top_aqi.index)
plt.title("Top 10 Countries by Average AQI (Historical)")
plt.xlabel("Average AQI (US EPA Index)")
plt.tight_layout()
plt.show()
```



```
[150]: # Compact correlation heatmap for selected features

selected_cols = [
    'temperature_celsius', 'humidity', 'pressure_mb',
    'wind_kph', 'uv_index', 'air_quality_PM2.5'
]

plt.figure(figsize=(8,6))
sns.heatmap(df[selected_cols].corr(), annot=True, cmap='viridis')
plt.title("Compact Correlation Heatmap: Key Features vs PM2.5")
plt.tight_layout()
plt.show()
```



Insights (Additional EDA)

- Wind speeds are mostly moderate; extremely high winds are rare.
- Humidity behavior changes across temperature bins; this helps interpret comfort and potential fog/haze conditions.
- AQI top-10 countries confirm and complement the PM2.5-based ranking.
- The compact correlation plot zooms into a smaller set of features important for PM2.5 modeling.

1.4 5. Machine Learning Setup

We now prepare features for ML and evaluate

1.4.1 Targets

- `air_quality_PM2.5` – fine particulate pollution
- `air_quality_us-epa-index` – air-quality index
- `temperature_celsius` – temperature in °C

We use **20+ features** (in practice, >30) from weather, air-quality precursors, geographic, and temporal information.

```
[151]: # 5.1 Define target columns
TARGET_PM25 = 'air_quality_PM2.5'
TARGET_AQI = 'air_quality_us-epa-index'
TARGET_TEMP = 'temperature_celsius'

# 5.2 Build feature DataFrame by dropping targets and raw datetime columns
drop_cols = [
    TARGET_PM25, TARGET_AQI, TARGET_TEMP,
    'last_updated', 'sunrise', 'sunset', 'moonrise', 'moonset',
    'last_updated_dt' # helper column from EDA
]

feature_df = df.drop(columns=drop_cols, errors='ignore').copy()

numeric_features = feature_df.select_dtypes(include=[np.number]).columns.
    ↪tolist()
categorical_features = feature_df.select_dtypes(include=['object']).columns.
    ↪tolist()

print("Number of numeric features:", len(numeric_features))
print("Number of categorical features:", len(categorical_features))
print("Total feature count used for ML:", len(numeric_features) + ↪
    len(categorical_features))
```

Number of numeric features: 39
 Number of categorical features: 4
 Total feature count used for ML: 43

We are using **well over 20 features** in the ML pipeline, satisfying the project requirement for a non-trivial, high-dimensional problem.

1.4.2 5.3 Load Models

```
[152]: # These .pkl files are created in the training notebook and simply loaded here.
# Make sure they are uploaded to the same working directory when running this ↪
notebook.

model_pm25 = joblib.load("best_model_pm2.5.pkl")
model_aqi = joblib.load("best_model_aqi.pkl")
model_temp = joblib.load("best_model_temperature.pkl")

print("Pre-trained models loaded successfully.")
```

Pre-trained models loaded successfully.

1.4.3 5.4 Evaluate Models and Create Prediction Columns

We evaluate each model on the full dataset (for demonstration) and also store predictions in new columns:

- pred_PM25
- pred_AQI
- pred_Temp

```
[153]: def evaluate_and_predict(model, X, y_true, col_name):
    y_pred = model.predict(X)
    mae = mean_absolute_error(y_true, y_pred)
    rmse = np.sqrt(mean_squared_error(y_true, y_pred))
    r2 = r2_score(y_true, y_pred)
    print(f'{col_name} - MAE: {mae:.3f}, RMSE: {rmse:.3f}, R2: {r2:.3f}')
    return y_pred, mae, rmse, r2

pred_pm25, mae_pm, rmse_pm, r2_pm = evaluate_and_predict(model_pm25, df[feature_df[TARGET_PM25], "PM2.5"])
pred_aqi, mae_aqi, rmse_aqi, r2_aqi = evaluate_and_predict(model_aqi, df[feature_df[TARGET_AQI], "AQI"])
pred_temp, mae_temp, rmse_temp, r2_temp = evaluate_and_predict(model_temp, df[feature_df[TARGET_TEMP], "Temperature"])

df['pred_PM25'] = pred_pm25
df['pred_AQI'] = pred_aqi
df['pred_Temp'] = pred_temp
```

PM2.5 - MAE: 0.844, RMSE: 2.779, R²: 0.995
AQI - MAE: 0.036, RMSE: 0.104, R²: 0.988
Temperature - MAE: 0.003, RMSE: 0.011, R²: 1.000

```
[154]: results_summary = pd.DataFrame({
    "Target": ["PM2.5", "AQI", "Temperature"],
    "MAE": [mae_pm, mae_aqi, mae_temp],
    "RMSE": [rmse_pm, rmse_aqi, rmse_temp],
    "R2": [r2_pm, r2_aqi, r2_temp]
})
results_summary
```

	Target	MAE	RMSE	R2
0	PM2.5	0.843615	2.778723	0.994975
1	AQI	0.036344	0.103571	0.988428
2	Temperature	0.002689	0.010845	0.999999

Insights (Model Performance)


- Models achieve reasonable RMSE and R² given the noisy, real-world nature of environmental data.
- Performance is typically best for temperature, then PM2.5, then AQI (since AQI aggregates multiple pollutants).
- These models are good enough for ranking and scenario analysis, which is the main goal of this project.


1.4.4 5.5 ML Visualizations: Actual vs Predicted, Error Distributions, and Feature Importance


```
[155]: def plot_pred_vs_actual(y_true, y_pred, title):
    plt.figure(figsize=(7,7))
    plt.scatter(y_true, y_pred, alpha=0.3)
    min_val = min(y_true.min(), y_pred.min())
    max_val = max(y_true.max(), y_pred.max())
    plt.plot([min_val, max_val], [min_val, max_val], 'r--')
    plt.xlabel("Actual")
    plt.ylabel("Predicted")
    plt.title(title)
    plt.grid(alpha=0.3)
    plt.tight_layout()
    plt.show()

plot_pred_vs_actual(df[TARGET_PM25], pred_pm25, "Actual vs Predicted PM2.5")
plot_pred_vs_actual(df[TARGET_AQI], pred_aqi, "Actual vs Predicted AQI")
plot_pred_vs_actual(df[TARGET_TEMP], pred_temp, "Actual vs Predicted
    ↵Temperature (°C)")
```

Actual vs Predicted PM2.5










```
[156]: # Error distributions

def plot_error_hist(errors, title):
    plt.figure(figsize=(8,5))
    sns.histplot(errors, bins=40, kde=True)
    plt.title(title)
    plt.xlabel("Prediction Error")
    plt.tight_layout()
    plt.show()

plot_error_hist(pred_pm25 - df[TARGET_PM25], "PM2.5 Prediction Error Distribution")
plot_error_hist(pred_aqi - df[TARGET_AQI], "AQI Prediction Error Distribution")
```

```
plot_error_hist(pred_temp - df[TARGET_TEMP], "Temperature Prediction Error Distribution")
```





```
[157]: # Feature importance for PM2.5 model, if supported (e.g., RandomForest / XGBoost)

try:
    # Pipeline structure: preprocessor + model
    model_step = model_pm25.named_steps.get('model', None)
    preproc = model_pm25.named_steps.get('preprocessor', None)

    if model_step is not None and hasattr(model_step, 'feature_importances_') and preproc is not None:
        # Numeric feature names from preprocessor
        num_feats = preproc.transformers_[0][2]
        # Categorical features will be one-hot encoded, but for simplicity we just show numeric here
        importances = model_step.feature_importances_[:len(num_feats)]


        sorted_idx = np.argsort(importances)[::-1][:15]
        top_feats = np.array(num_feats)[sorted_idx]
        top_imps = importances[sorted_idx]

        plt.figure(figsize=(10,6))
        sns.barplot(x=top_imps, y=top_feats)
        plt.title("Top 15 Numeric Feature Importances for PM2.5 Model")
        plt.xlabel("Importance")
```

```

        plt.tight_layout()
        plt.show()
    else:
        print("Model does not expose feature_importances_; skipping feature_importance plot.")
except Exception as e:
    print("Error while computing feature importances:", e)

```



Insights (ML Visuals)

- Scatter plots show how close predictions are to the 45° line (perfect agreement).
- Error histograms help us see whether models are biased (e.g., always under-predicting high values).
- Feature importance (for tree-based models) highlights which numeric weather features matter most for predicting PM2.5 (e.g., humidity, wind, pressure).
- These visualizations go beyond simple metrics and help explain **how** and **why** the models behave as they do.


1.4.5 5.6 Prediction-Only Visualizations (How Model Output Looks)

Here we visualize **only the model predictions**, independent of the historical values.

Example: Top 10 countries by **predicted** PM2.5 (current conditions).

```
[158]: pred_pm25_country = df.groupby('country')['pred_PM25'].mean().
    sort_values(ascending=False).head(10)

plt.figure(figsize=(10,6))
sns.barplot(x=pred_pm25_country.values, y=pred_pm25_country.index)
plt.title("Top 10 Countries by Predicted PM2.5 (Current Scenario)")
plt.xlabel("Predicted Avg PM2.5")
plt.tight_layout()
plt.show()
```



Insight (Prediction-Only View)

This plot answers the question: “*If we trust the ML model’s understanding of relationships, which countries look worst under current-like conditions?*”

It is particularly useful for **communicating model output to non-technical stakeholders** who just want a ranking or risk list.

1.5 6. Scenario-Based Predictive Country Ranking (Future-like Conditions)

The dataset is a **single-day snapshot**, not a historical time-series.

Therefore, we do **scenario forecasting** rather than strict date-based forecasting.

We simulate a **future-like scenario** by shifting the hour-of-day feature (`last_updated_hour`) forward and using the models to recompute predictions.

This allows us to estimate **which countries are likely to remain or become high-risk** under slightly changed environmental conditions.

```
[159]: # 6.1 Build future-like feature DataFrame

future_df = feature_df.copy()

# Use current hour/minute from cleaned df
future_df['last_updated_hour'] = df['last_updated_hour']
future_df['last_updated_minute'] = df['last_updated_minute']

# Shift hour by +6 (wrap around 24 hours)
future_df['future_hour_shifted'] = (df['last_updated_hour'] + 6) % 24
future_df['last_updated_hour'] = future_df['future_hour_shifted']
future_df.drop(columns=['future_hour_shifted'], inplace=True)

# 6.2 Predict under the future-like scenario
future_df['pred_PM25'] = model_pm25.predict(future_df)
future_df['pred_AQI'] = model_aqi.predict(future_df)
future_df['pred_Temp'] = model_temp.predict(future_df)
```

```
[160]: # 6.3 Aggregate country-level predictions for the scenario

pm25_country_future = (
    future_df
    .join(df[['country']], rsuffix='_orig')
    .groupby('country')['pred_PM25']
    .mean()
)

aqi_country_future = (
    future_df
    .join(df[['country']], rsuffix='_orig')
    .groupby('country')['pred_AQI']
    .mean()
)

temp_country_future = (
    future_df
    .join(df[['country']], rsuffix='_orig')
    .groupby('country')['pred_Temp']
    .mean()
)

def top_bottom(series, n=10):
    return (
        series.sort_values(ascending=False).head(n),
        series.sort_values(ascending=True).head(n)
    )
```

```

pm25_top_future, pm25_bottom_future = top_bottom(pm25_country_future)
aqi_top_future, aqi_bottom_future = top_bottom(aqi_country_future)
temp_top_future, temp_bottom_future = top_bottom(temp_country_future)

print("== Predicted Top 10 PM2.5 Countries (Future-like Scenario) ==")
display(pm25_top_future.to_frame("Predicted_PM2.5"))

```

== Predicted Top 10 PM2.5 Countries (Future-like Scenario) ==

Predicted_PM2.5

country	
Chile	177.531602
Saudi Arabia	140.773086
China	136.601890
India	111.545970
Kuwait	97.472221
Indonesia	93.517034
Bahrain	71.645213
Mauritania	71.097183
Südkorea	69.406692
Bangladesh	68.277793

[161]: # Visualize future-scenario top/bottom 10 PM2.5

```


fig, axes = plt.subplots(1, 2, figsize=(18,6))

sns.barplot(x=pm25_top_future.values, y=pm25_top_future.index, ax=axes[0])
axes[0].set_title("Future Scenario - Top 10 Predicted PM2.5 Countries")
axes[0].set_xlabel("Predicted PM2.5")

sns.barplot(x=pm25_bottom_future.values, y=pm25_bottom_future.index, ax=axes[1])
axes[1].set_title("Future Scenario - Bottom 10 Predicted PM2.5 Countries")
axes[1].set_xlabel("Predicted PM2.5")

plt.tight_layout()
plt.show()

```



Insights (Scenario Forecasting)

- Countries consistently appearing in both historical and predicted top-10 lists are clear **hotspots for intervention**.
- Countries moving into or out of the predicted top-10 may signal **changing risk profiles** under different time-of-day or environmental scenarios.
- This approach is transparent about limitations: it is scenario-based, not a multi-year climate forecast.

1.6 7. Exporting Data for Dashboards

We export:

1. A **full dataset** including all original variables, engineered features, and prediction columns.
2. A **single compact CSV** that holds **all top/bottom 10 historical and predicted rankings** for PM2.5, AQI, and Temperature.

```
[162]: # 7.1 Final dataset with predictions (for Tableau / Power BI)
```

```
final_export_df = df.copy()

# Add scenario predictions from future_df (aligned by index)
final_export_df['scenario_pred_PM25'] = future_df['pred_PM25']
final_export_df['scenario_pred_AQI'] = future_df['pred_AQI']
final_export_df['scenario_pred_Temp'] = future_df['pred_Temp']

export_path_full = "Global_Weather_Final_With_Predictions.csv"
final_export_df.to_csv(export_path_full, index=False)

print("Full dataset with predictions saved to:")
print(export_path_full)
```

Full dataset with predictions saved to:
Global_Weather_Final_With_Predictions.csv

```
[163]: # 7.2 Combined Top/Bottom 10 historical + scenario rankings into ONE CSV
```

```
combined_rows = []

def add_block(series, metric_name, category_name):
    df_block = series.to_frame(name="value").reset_index()
    df_block["metric"] = metric_name
    df_block["category"] = category_name
    return df_block

# Historical
combined_rows.append(add_block(pm25_top_hist, "PM2.5", "Historical Top 10"))
```

```

combined_rows.append(add_block(pm25_bottom_hist, "PM2.5", "Historical Bottom\u219210"))
combined_rows.append(add_block(aqi_top_hist, "AQI", "Historical Top 10"))
combined_rows.append(add_block(aqi_bottom_hist, "AQI", "Historical Bottom 10"))
combined_rows.append(add_block(temp_top_hist, "Temperature", "Historical Top\u219210"))
combined_rows.append(add_block(temp_bottom_hist, "Temperature", "Historical\u2192Bottom 10"))

# Scenario
combined_rows.append(add_block(pm25_top_future, "PM2.5", "Scenario Top 10"))
combined_rows.append(add_block(pm25_bottom_future, "PM2.5", "Scenario Bottom\u219210"))
combined_rows.append(add_block(aqi_top_future, "AQI", "Scenario Top 10"))
combined_rows.append(add_block(aqi_bottom_future, "AQI", "Scenario Bottom 10"))
combined_rows.append(add_block(temp_top_future, "Temperature", "Scenario Top\u219210"))
combined_rows.append(add_block(temp_bottom_future, "Temperature", "Scenario\u2192Bottom 10"))

combined_df = pd.concat(combined_rows, ignore_index=True)
export_path_rankings = "TopBottom_Rankings_Historical_vs_Scenario.csv"
combined_df.to_csv(export_path_rankings, index=False)

print("Combined rankings CSV saved to:")
print(export_path_rankings)

combined_df.head()

```

Combined rankings CSV saved to:
TopBottom_Rankings_Historical_vs_Scenario.csv

[163]:

	country	value	metric	category
0	Chile	178.947781	PM2.5	Historical Top 10
1	Saudi Arabia	140.220979	PM2.5	Historical Top 10
2	China	137.501838	PM2.5	Historical Top 10
3	India	110.518835	PM2.5	Historical Top 10
4	Kuwait	98.752979	PM2.5	Historical Top 10

1.7 8. Project Summary & Connection to Real-World Goals

- The dataset is **large, messy, and rich**, with 40+ original features and additional engineered time and categorical encodings.
- Cleaning and feature engineering created a consistent foundation for both **EDA** and **ML**.
- EDA revealed global patterns in weather and pollution, including **country-level hotspots** and **time-of-day effects**.

- ML models (Random Forest / XGBoost / KNN chosen in a separate training notebook) were **loaded as pre-trained artifacts** and evaluated here using multiple metrics and visualizations.
- Scenario-based forecasting used these models to generate **future-like rankings of countries** by predicted PM2.5, AQI, and temperature.
- Final enriched datasets and ranking tables are exported for **interactive dashboards**, satisfying the course requirements for integrated analytics and visualization.

This notebook can directly support the **technical report**, **final presentation**, and **dashboard demo**.